

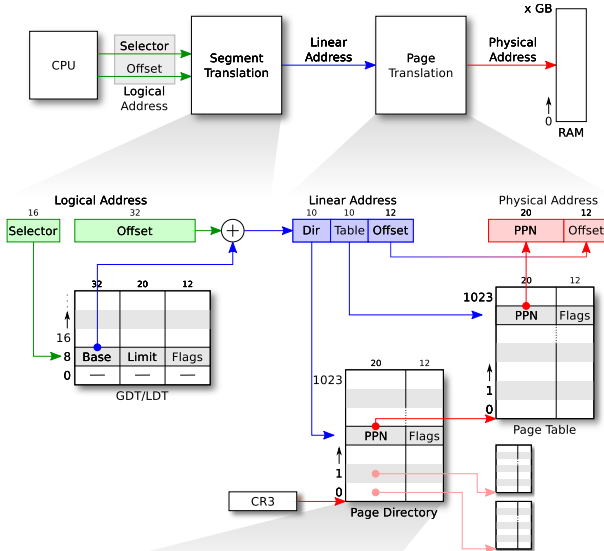
236366 Operating Systems Engineering

Recitation #3 (part 1):

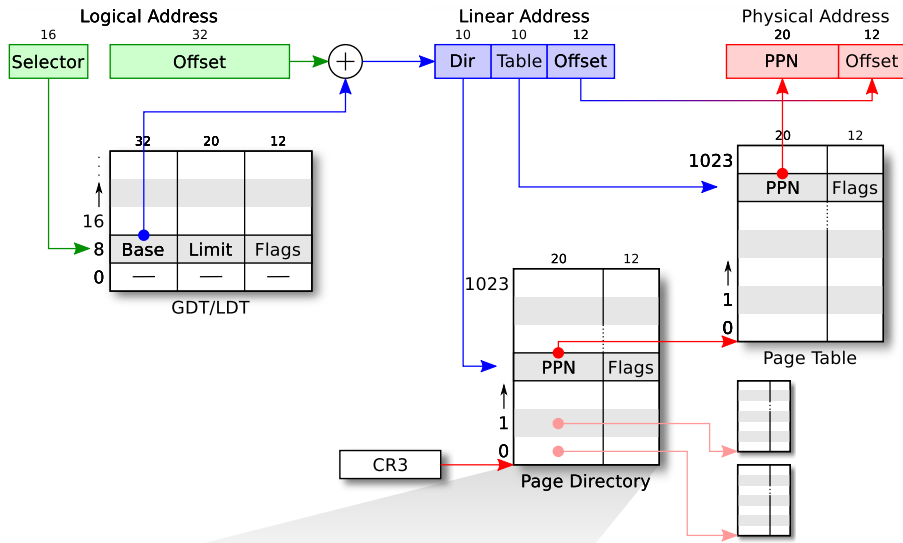
Address Translation and Paging

Presented by:
Ilia Kravets
<ilia@cs>



x86 address translation



x86 address translation



x86 address translation



Page table and page directory entries are identical except for the D bit.

- P** - Present
- W** - Writable
- U** - User
- WT** - 1=Write-through, 0=Write-back
- CD** - Cache Disabled
- A** - Accessed
- D** - Dirty (0 in page directory)
- AVL** - Available for system use

Using x86 paging

- start using paging
 - setup PD, PT's
 - point to PD: `%cr3 = PHYSADDR(PD)`
 - enable paging: `%cr0 = %cr0 | CR0_PG`
 - may “disable” segmentation
- flush TLB when updating PD/PT
 - full flush: reload `%cr0`
 - specific entry only: `invlpg VA`
- access control
 - S/U bit
 - R/W bit (& `CR0_WP`)
 - `%cr3` can be modified by kernel only
 - what about PD/PT?

Using x86 paging

- start using paging
 - setup PD, PT's
 - point to PD: `%cr3 = PHYSADDR(PD)`
 - enable paging: `%cr0 = %cr0 | CR0_PG`
 - may “disable” segmentation
- flush TLB when updating PD/PT
 - full flush: reload `%cr0`
 - specific entry only: `invlpg VA`
- access control
 - S/U bit
 - R/W bit (& `CR0_WP`)
 - `%cr3` can be modified by kernel only
 - what about PD/PT?

Using x86 paging

- start using paging
 - setup PD, PT's
 - point to PD: `%cr3 = PHYSADDR(PD)`
 - enable paging: `%cr0 = %cr0 | CR0_PG`
 - may “disable” segmentation
- flush TLB when updating PD/PT
 - full flush: reload `%cr0`
 - specific entry only: `invlpg VA`
- access control
 - S/U bit
 - R/W bit (& `CR0_WP`)
 - `%cr3` can be modified by kernel only
 - what about PD/PT?

MMU: linear to physical address translation

```
uint translate (uint la, bool user, bool write)
{
    uint pde;
    pde = read_mem (%CR3 + 4*(la >> 22));
    access (pde, user, write);
    pte = read_mem ( (pde & 0xfffff000) + 4*((la >> 12) & 0x3ff));
    access (pte, user, write);
    return (pte & 0xfffff000) + (la & 0xfff);
}

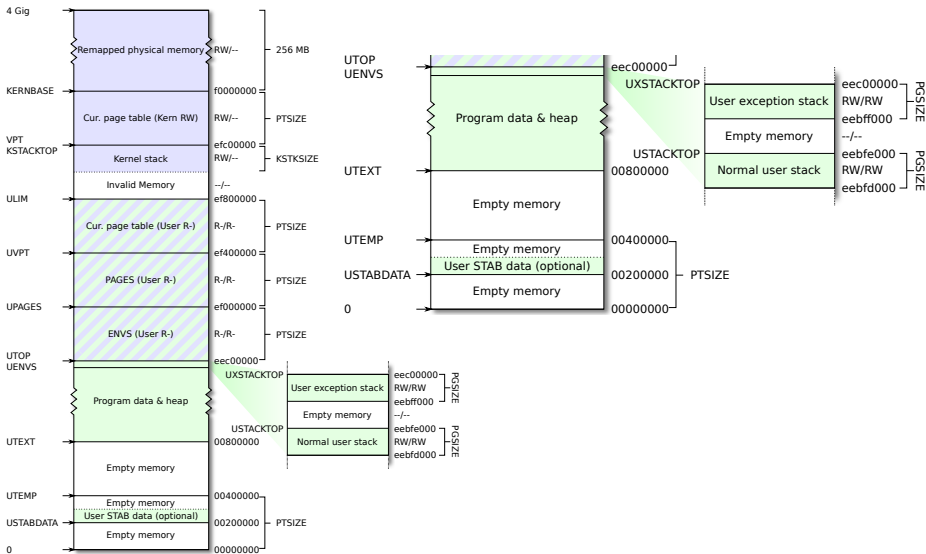
// check protection. pxe is a pte or pde.
// user is true if CPL==3
void access (uint pxe, bool user, bool write)
{
    if (!(pxe & PG_P)
        => page fault — page not present
    if (!(pxe & PG_U) && user)
        => page fault — not access for user

    if (write && !(pxe & PG_W)) {
        if (user)
            => page fault — not writable
        if (%CR0 & CR0_WP)
            => page fault — not writable
    }
}
```

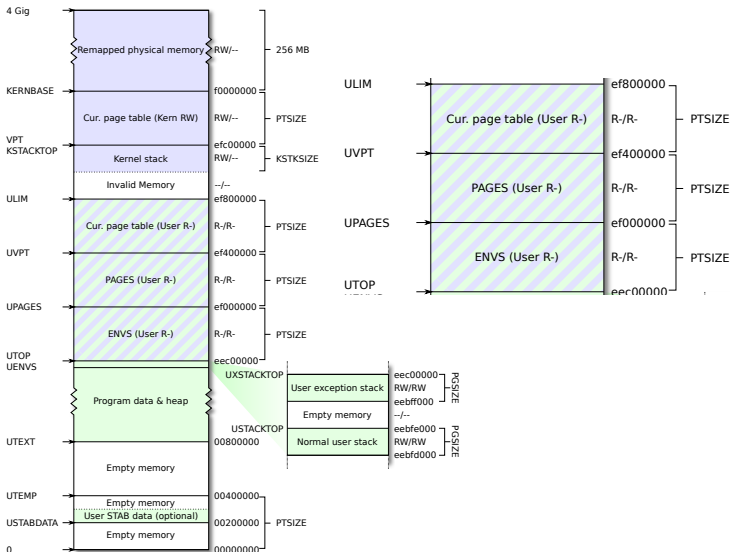

JOS virtual address space

- segmentation is “disabled”
 - identity $VA \rightarrow LA$ mapping (base=0, limit=4G)
 - user/kernel differs in DPL only (about it later. . .)
- per process page table
- same address space for both user and kernel
 - kernel is at the top
 - kernel part is the same for all processes
- limit user access (non-user addresses are usually inaccessible)
- JOS tricks
 - some kernel memory is remapped read only for user
 - physical memory is remapped continuously
 - virtual page table (VPT)

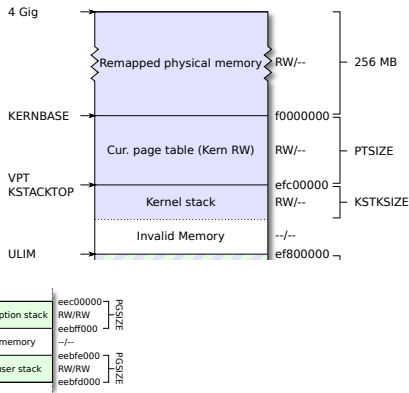
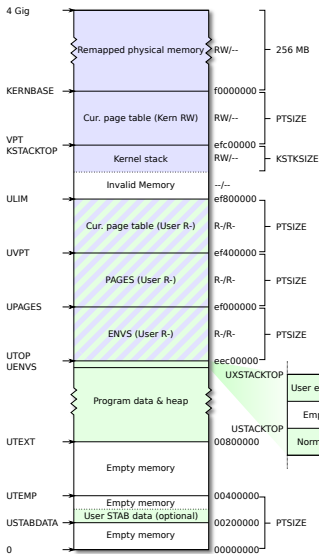
JOS virtual address space map



JOS virtual address space map



JOS virtual address space map



JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - how and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - we have "self-maintained" mapping
 - put physical address of PD into PD's entry # x
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - *that's* what we need to do when a PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - *that's* how we do "self-modifying" mapping
 - put physical address of PD into PD's entry # x
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - *that's* how to setup!

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - *that's* what we call a *second-level PT*
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - we have "second-level PT" mapping
 - put physical address of PD into PD's entry # x
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - *2nd level PTs are not updated*
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - *2nd level PTs are not updated*
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - *2nd level PTs are not updated*
 - UVPT – read-only user's view of VPT
 - *2nd level PTs are not updated*

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
- UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?

JOS virtual page table

- PT is just a mapping function $f(\text{linear}) \rightarrow \text{physical}$
- implemented in x86 as 2-level page walk
- OS needs to re-implement $f()$ in software
- wouldn't it be nice to just use `pagetable[linear]`?
- so 2nd level PTs should form a continuous 4MB array
 - continuous in *virtual* memory
 - MMU translates all memory accesses using PT, remember?
 - just need to setup one 2nd level PT which will map *all* 2nd level PTs
 - and maintain it whenever PD is updated
 - but PD does exactly this!
 - let's just use PD as 2nd level PT
 - bonus: "self-maintained" mapping
 - put physical address of PD into PD's entry $\#x$
 - we got a 4MB array starting at VA $x \ll 22$
 - what's the value of x in JOS?
 - UVPT – read-only user's view of VPT
 - how to setup?